



清华大学
Tsinghua University



MIT



Twilight: Adaptive Attention Sparsity with Hierarchical Top- p Pruning

**Chaofan Lin, Jiaming Tang, Shuo Yang, Hanshuo Wang, Tian Tang, Boyu Tian, Ion Stoica, Song Han,
Mingyu Gao**

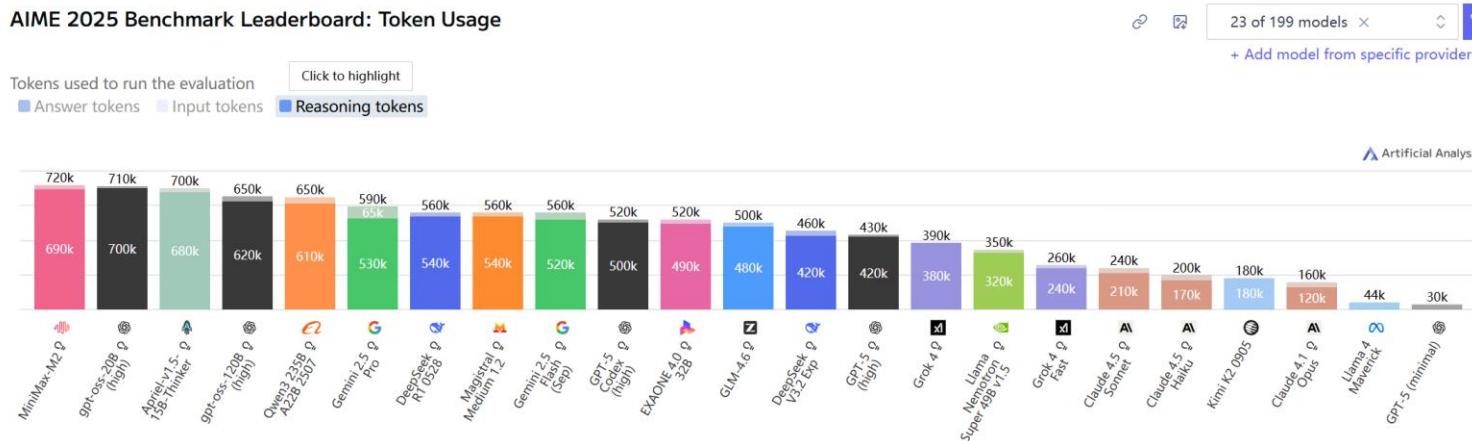
Tsinghua University
Massachusetts Institute of Technology
University of California, Berkeley

<https://github.com/tsinghua-ideal/Twilight>


Background

Long-context LLMs are powerful but computationally expensive

- ❑ Trend: **Long Context** windows are becoming the new standard for state-of-the-art LLMs, especially for reasoning models.
- ❑ For long context LLM inference, **attention** dominates the latency.



Reasoning tasks like AIME cost nearly 1M tokens




Attention becomes the bottleneck as the sequence length increases


Background

Top- k Sparse Attention to reduce KV cache loading

- Since attention is **memory-bound**, previous works propose sparse attention, which first estimates attention scores then selectively loads only important tokens.
- However, the main challenge of top- k sparse attention is to find a **universally applicable budget** to all scenarios.



Dynamism Across Prompts (Budget/Prompt Len)	Prompt 0 158/10739 =1.47%	Prompt 1 408/17501 =2.33%	Prompt 2 244/5441 =4.48%	Prompt 3 195/2023 =9.64%
Dynamism Across Queries (Budget)	Query 0 120	Query 1 152	Query 2 174	Query 3 185
Dynamism Across Layers (Budget)	Layer 6 27	Layer 10 356	Layer 15 105	Layer 26 228
Dynamism Across Heads (Budget)	Head 0 37	Head 6 2707	Head 15 886	Head 23 144



Previous work (Quest)

The best budget choices vary dynamically across different levels.


Bringing Top- p Sampling to Sparse Attention

Top- p Sparse Attention is inherently budget-adaptive


- We argue that the core reason for budget dynamism is the dynamic nature of the attention weight distributions at runtime, thus propose Top- p Sparse Attention.

Definition 3.3 (Oracle Top- p Sparse Attention). Given the threshold p ,

$$\mathcal{I} = \arg \min_{\mathcal{I}} |\mathcal{I}| \quad \text{s.t. } \sum_{i \in \mathcal{I}} \mathbf{W}[i] \geq p$$



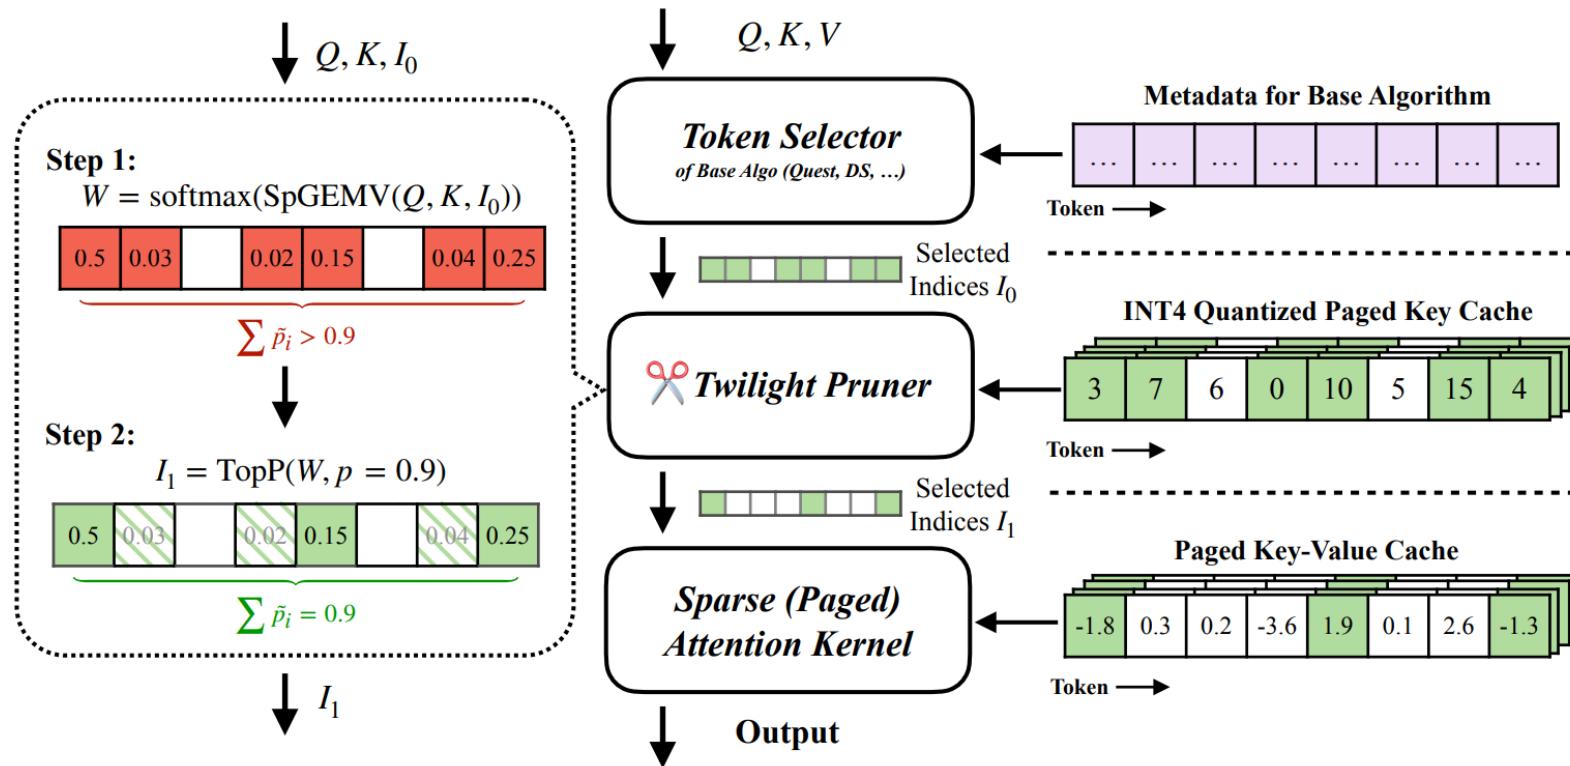
Diverse distributions observed in attention weights of different attention heads.



Cumulative attention scores of different budget selections in one example attention head.

Twilight


- Key Design: Hierarchical **Select-then-Prune** architecture as a unified optimizer for all existing top- k based sparse attention methods (denoted as *BaseAlgo*).
- First *BaseAlgo* uses a conservative, relatively large budget. Then **Twilight** further prunes them using efficient top- p Pruner.



Twilight

To Achieve the Efficient Pruner


- Efficient SpGEMV with 4-bit Quantization of Key Cache to estimate token importance: we find that 4-bit strikes a balance between accuracy and efficiency.
- Efficient Sorting-free Top-p via binary search modified from FlashInfer.
- Load Balancing with Awareness of Head Dynamism with GQA adaption.



Algorithm 1 Top- p via Binary Search.

Input: normalized attention weights $W \in \mathbb{R}^{BS \times H \times N}$, top- p threshold p , hyper-parameter ϵ .
Output: indices \mathcal{I} , mask $\mathcal{M} \in \{0, 1\}^{BS \times H \times N}$.

```
 $l = 0, r = \max(W), m = (l + r)/2;$ 
repeat
   $W_0 = \text{where}(W < m, 0.0, W);$ 
   $W_1 = \text{where}(W \leq l, \text{INF}, W);$ 
   $W_2 = \text{where}(W > r, -\text{INF}, W);$ 
  if  $\text{sum}(W_0) \geq p$  then
     $l = m;$ 
  else
     $r = m;$ 
  end if
until  $\max(W_2) - \min(W_1) \geq \epsilon$ 
Select indices  $\mathcal{I}$  and set mask  $\mathcal{M}$  where  $W \geq l$ ;
return  $\mathcal{I}, \mathcal{M}$ ;
```



Accuracy Evaluation

- Twilight achieves nearly **no accuracy loss** on three medium-context benchmarks and two long-context benchmarks (LongBench, RULER).

Table 2: Average scores on 12 different tasks from Longbench. We report relative error changes (**improvement** or **degradation**) when integrating Twilight with each base algorithm. Detailed results are in [Table 5](#) in [Appendix C](#).

	Budget	Longchat-7B -v1.5-32k	LLaMA-3.1-8B -Instruct
Full	32k	36.78	52.01
	Twilight	38.52 (+4.7%)	51.64 (-0.7%)
MagicPIG	K=8, L=75	-	51.70
	K=10, L=150	-	51.32
Quest	256	31.26	38.20
	1024	36.85	47.79
	4096	37.33	50.79
	8192	37.10	51.44
	Twilight	38.04 (+2.5%)	51.57 (+0.3%)
DS	256	35.32	45.74
	1024	35.96	49.43
	4096	36.31	50.98
	8192	36.62	51.14
	Twilight	38.71 (+5.7%)	51.73 (+1.2%)

Table 3: Average scores on RULER.

	Budget	16k	32k	64k	96k	Avg.
Full	100%	92.88	89.42	85.17	85.23	88.18
	Twilight	93.13	89.10	84.64	83.10	87.49
MagicPIG	K=8, L=75	92.22	89.37	84.07	82.58	87.06
	K=10, L=150	91.38	88.20	83.34	82.02	86.23
Quest	4%	79.35	79.8	78.64	73.22	77.75
	8%	87.31	83.06	80.82	75.28	81.62
	Twilight	91.53	87.97	84.12	82.96	86.65
DS	4%	92.04	88.11	84.43	82.56	86.79
	8%	92.89	88.70	84.39	82.72	87.18
	Twilight	93.54	89.24	85.91	82.81	87.88

Table 4: Results on 3 medium-context benchmarks.

	GSM8K(flexible/strict)↑ COQA(em/f1)↑ PG-19 Perplexity↓	LLaMA-2-7B-Chat	LLaMA-3.1-8B-Instruct
Full	0.2290/0.2282	0.5935/0.7511	7.503
Quest	0.0523/0.0508	0.5710/0.7425	14.15
DS	0.2191/0.2190	0.5855/0.7401	7.622
Twilight	0.2153/0.2115	0.6088/0.7642	7.600
(Twilight Avg. Budget)	90.82	91.86	102.58
		LLaMA-3.1-8B-Instruct	
Full	0.7726/0.7475	0.6363/0.7882	7.490
Quest	0.3639/0.3533	0.6007/0.7554	19.00
DS	0.6194/0.6027	0.6455/0.7964	7.967
Twilight	0.7771/0.7604	0.6325/0.7869	7.529
(Twilight Avg. Budget)	112.40	86.85	110.98

Efficiency Evaluation

- Twilight accelerates self-attention operator by $2.4\times$ (FlashInfer) and $1.4\times$ (Quest) at batch size=64
- And for E2E per-token latency, Quest-Twi is $3.9\times$ compared to FlashInfer and $1.35\times$ to Quest at batch size=256.




Figure 7: Latencies and speedups of self-attention at different sequence lengths and batch sizes.




Figure 8: Time-Per-Output-Token (TPOT) improvements in end-to-end serving scenarios.

- ❑ We propose Twilight, a composable optimizer to accelerate any existing top- k sparse decoding methods through **hierarchical top- p pruning**, making them **efficient and budget-adaptive**.
- ❑ Paper: <https://arxiv.org/abs/2502.02770>
- ❑ Code: <https://github.com/tsinghua-ideal/Twilight>

Thanks for Listening

Chaofan Lin

lcf24@mails.tsinghua.edu.cn



清华大学
Tsinghua University